

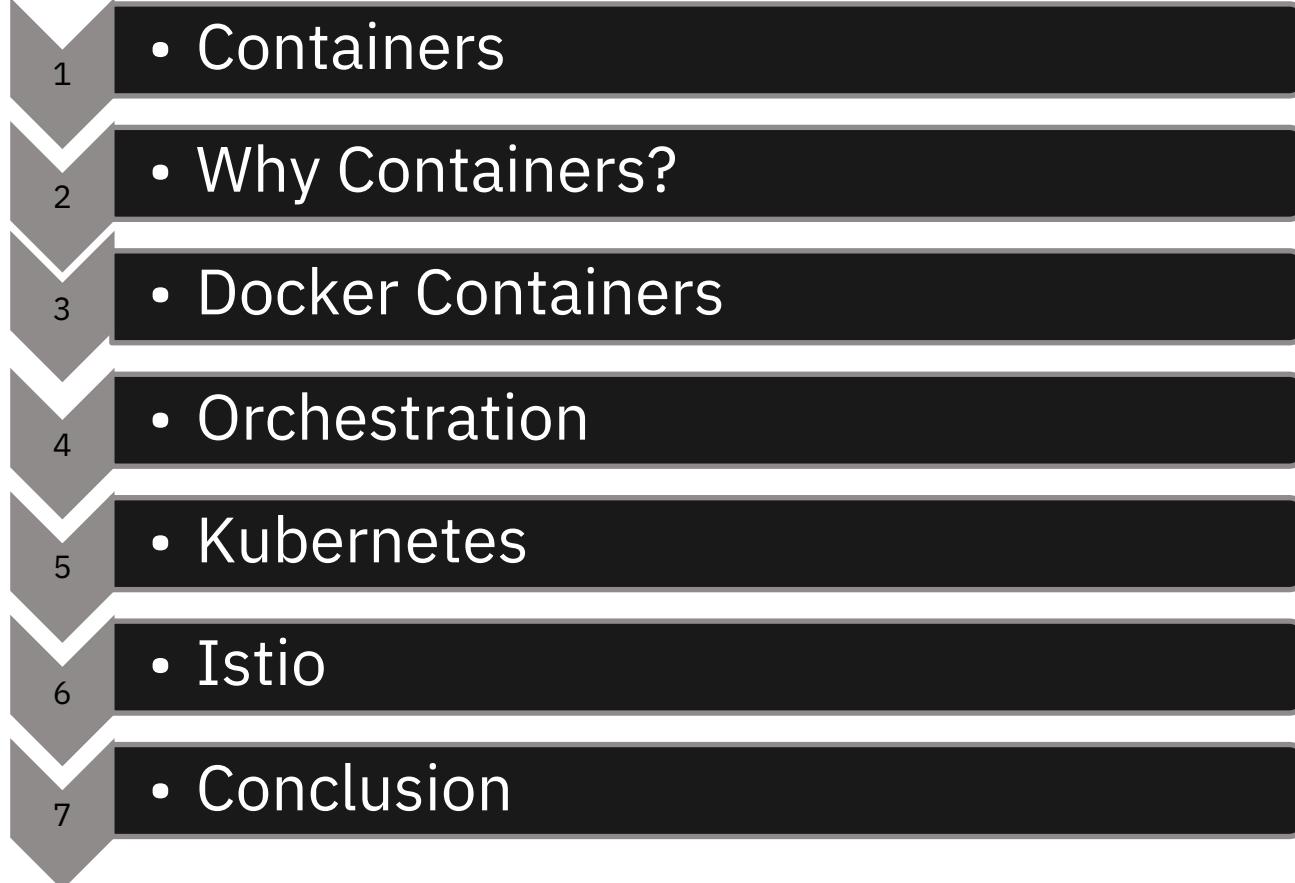
Serverless and Containers in Cloud Age – part two

IBM Code

Marek Sadowski
Developer Advocate | IBM San Francisco

@blumareks

Agenda

- 
- 1 • Containers
 - 2 • Why Containers?
 - 3 • Docker Containers
 - 4 • Orchestration
 - 5 • Kubernetes
 - 6 • Istio
 - 7 • Conclusion

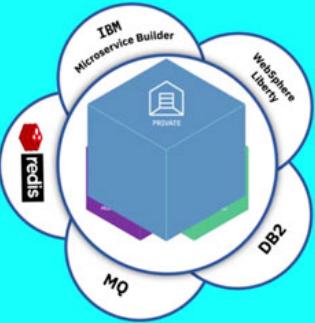
Containers

A standard way to package an application and all its dependencies so that it can be moved between environments and run without changes.

Containers work by isolating the differences between applications inside the container so that everything outside the container can be standardized.

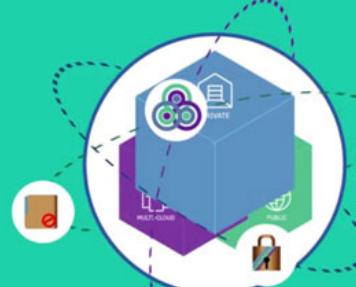
Why Containers?

Use Cases for Containers



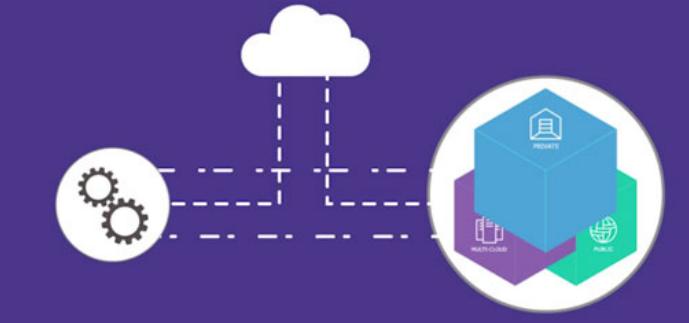
Create new cloud-native apps

Streamline development with built-in microservices, runtimes, containers and Kubernetes orchestration plus integrated management.



Modernize your heritage apps on cloud

Move your apps as-is to the cloud or refactor an app and use it in new development and application workload models.

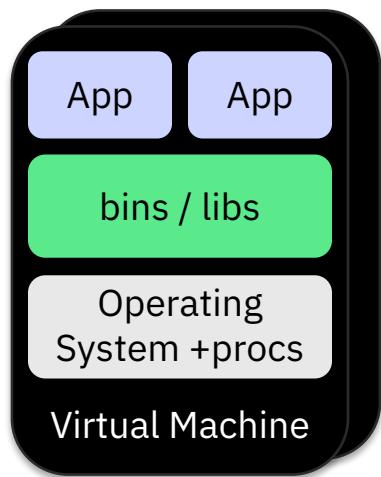


Open your data center to work with cloud services

Protect and leverage your in-house data and pull in external

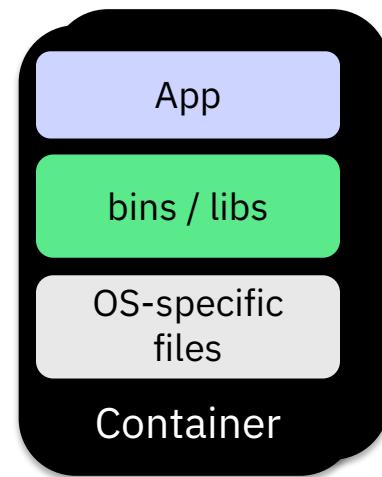
VM vs Container

Virtual Machine



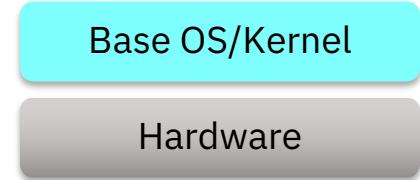
VS

Container



App, bins/libs/OS must all be runnable on the shared kernel

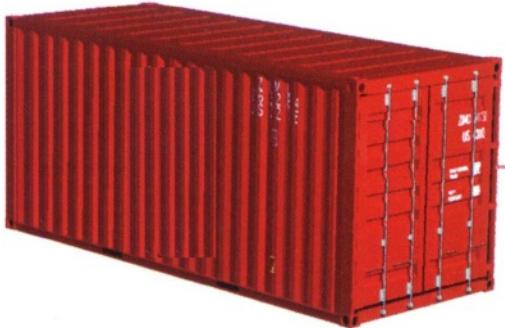
If OS files aren't needed they can be excluded.



Each VM has its own OS

Containers share the same base Kernel

Dev vs. Ops



- | Dev | Ops |
|--|--|
| <ul style="list-style-type: none">• Code• Libraries• Configuration• Server runtime• OS | <ul style="list-style-type: none">• Logging• Remote access• Network configuration• Monitoring |

Separation of concerns

A container separates and bridges the Dev and Ops in DevOps

- Dev focuses on the application environment
- Ops focuses on the deployment environment

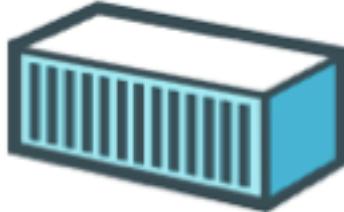
Docker Containers

Docker Mission

Docker is an **open platform** for building distributed applications for **developers** and **system administrators**



Build



Ship



Run



openstack™



Windows Azure



amazon
web services™



IBM Cloud

IBM Code



Google Compute Engine

Any App

Anywhere

What is Docker?

Tooling to manage containers

- Containers are not new
- Docker just made them easy to use

Docker creates and manages the lifecycle of containers

- Setup filesystem
- CRUD container
 - Setup networks
 - Setup volumes / mounts
 - Create: start new process telling OS to run it in isolation



Our First Container

```
$ docker run ubuntu echo Hello World  
Hello World
```

Our First Container

```
$ docker run ubuntu echo Hello World
```

```
Hello World
```

What happened?

Docker created a directory with a "ubuntu" filesystem (image)

Docker created a new set of namespaces

Ran a new process: echo Hello World

Using those namespaces to isolate it from other processes

Using that new directory as the "root" of the filesystem (chroot)

That's it!

Notice as a user I never installed "ubuntu"

Run it again - notice how quickly it ran

```
[Mareks-MBP:~ mareksadowski$ docker run ubuntu echo Hello World
Unable to find image 'ubuntu:latest' locally
latest: Pulling from library/ubuntu
6b98dfc16071: Pull complete
4001a1209541: Pull complete
6319fc68c576: Pull complete
b24603670dc3: Pull complete
97f170c87c6f: Pull complete
Digest: sha256:5f4bdc3467537cbbe563e80db2c3ec95d548a9145d64453b06
Status: Downloaded newer image for ubuntu:latest
Hello World
[Mareks-MBP:~ mareksadowski$ docker run ubuntu echo Hello World
Hello World
Mareks-MBP:~ mareksadowski$ ]
```

ssh-ing into a container - fake it...

```
$ docker run -ti ubuntu bash
```

Now the process is "bash" instead of "echo"

But its still just a process

Look around, mess around, its totally isolated

– MAKE SURE YOU'RE IN A CONTAINER!

```
Mareks-MBP:~ mareksadowski$ docker run -ti ubuntu bash
root@c6d6a09b9db9:/# pwd
/
root@c6d6a09b9db9:/# time

real    0m0.000s
user    0m0.000s
sys     0m0.000s
root@c6d6a09b9db9:/# du
12      ./root
3836    ./sbin
4       ./run/lock
4       ./run/mount
8       ./run/systemd
```

A look under the covers

```
$ docker run ubuntu ps -ef
UID          PID      PPID      C
STIME        TTY      TIME      CMD
root          1          0      0
14:33 ?      00:00:00  ps -ef
```

Things to notice with these examples

- Each container only sees its own process(es)
- Running as "root"
- Running as PID 1

Docker Component Overview

Docker Engine

- Manages containers on a host
- Accepts requests from clients
 - REST API
- Maps container ports to host ports
 - E.g. 80 → 3582

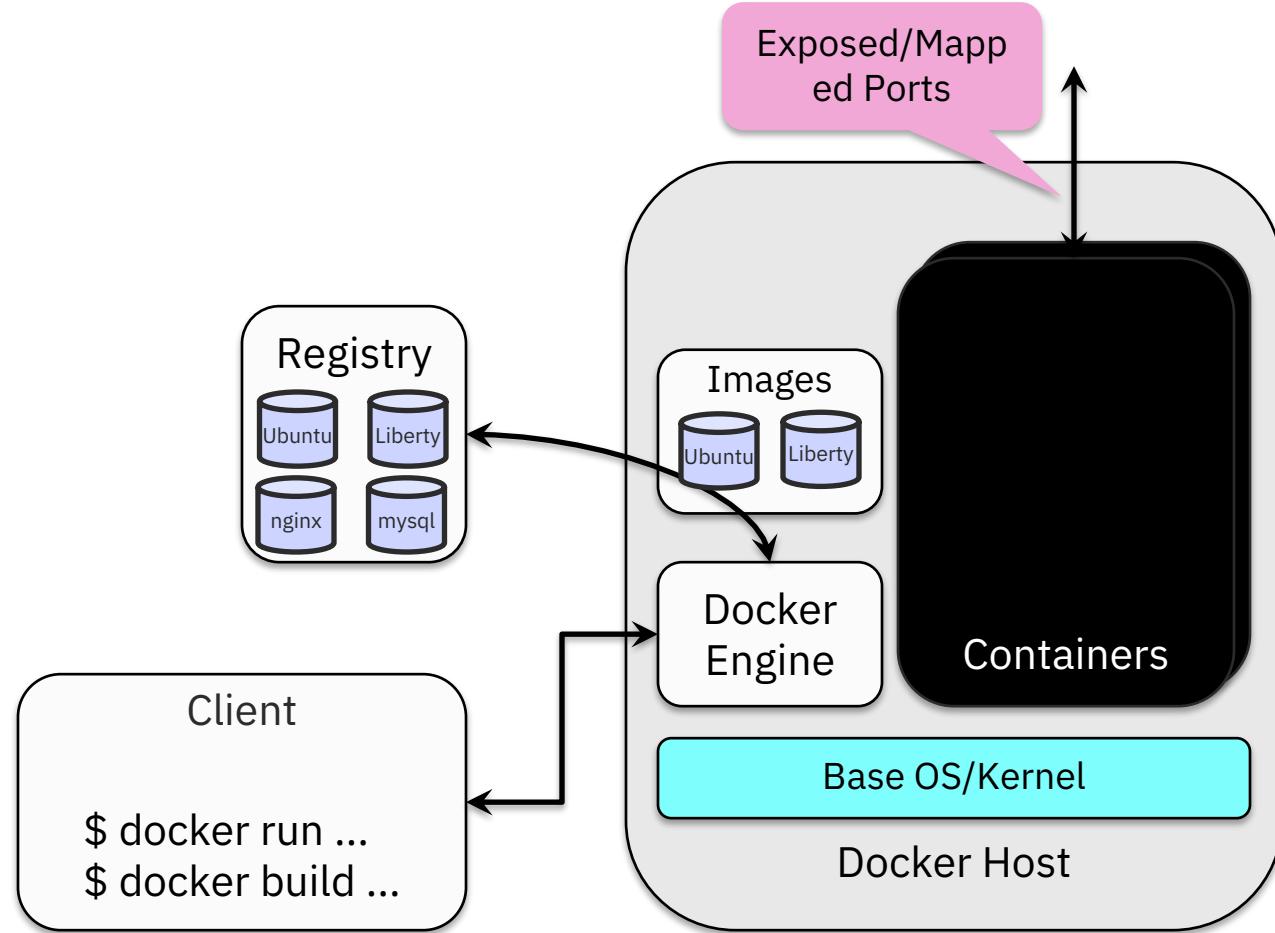
Images

Docker Client

- Drives engine
- Drives "builder" of Images

Docker Registry

- Image DB
IBM Code



IBM Cloud – Enterprise grade registry

IBM Cloud

Catalog Docs Support Manage

All Categories

Search Filter

Compute

Containers >

Networking

Storage

AI

Analytics

Databases

Developer Tools

Integration

Internet of Things

Security and Identity

Starter Kits

Web and Mobile

Application Services

Build your virtual environments.

 IBM Cloud Kubernetes Service

Deploy secure, highly available apps in a native Kubernetes experience.

IBM

 Container Registry

Store and distribute Docker container images with this fully managed private registry. M

Lite IBM

FEEDBACK

Containers



Everyone's container journey starts with one container....

Containers



At first the growth is easy to handle....



Containers



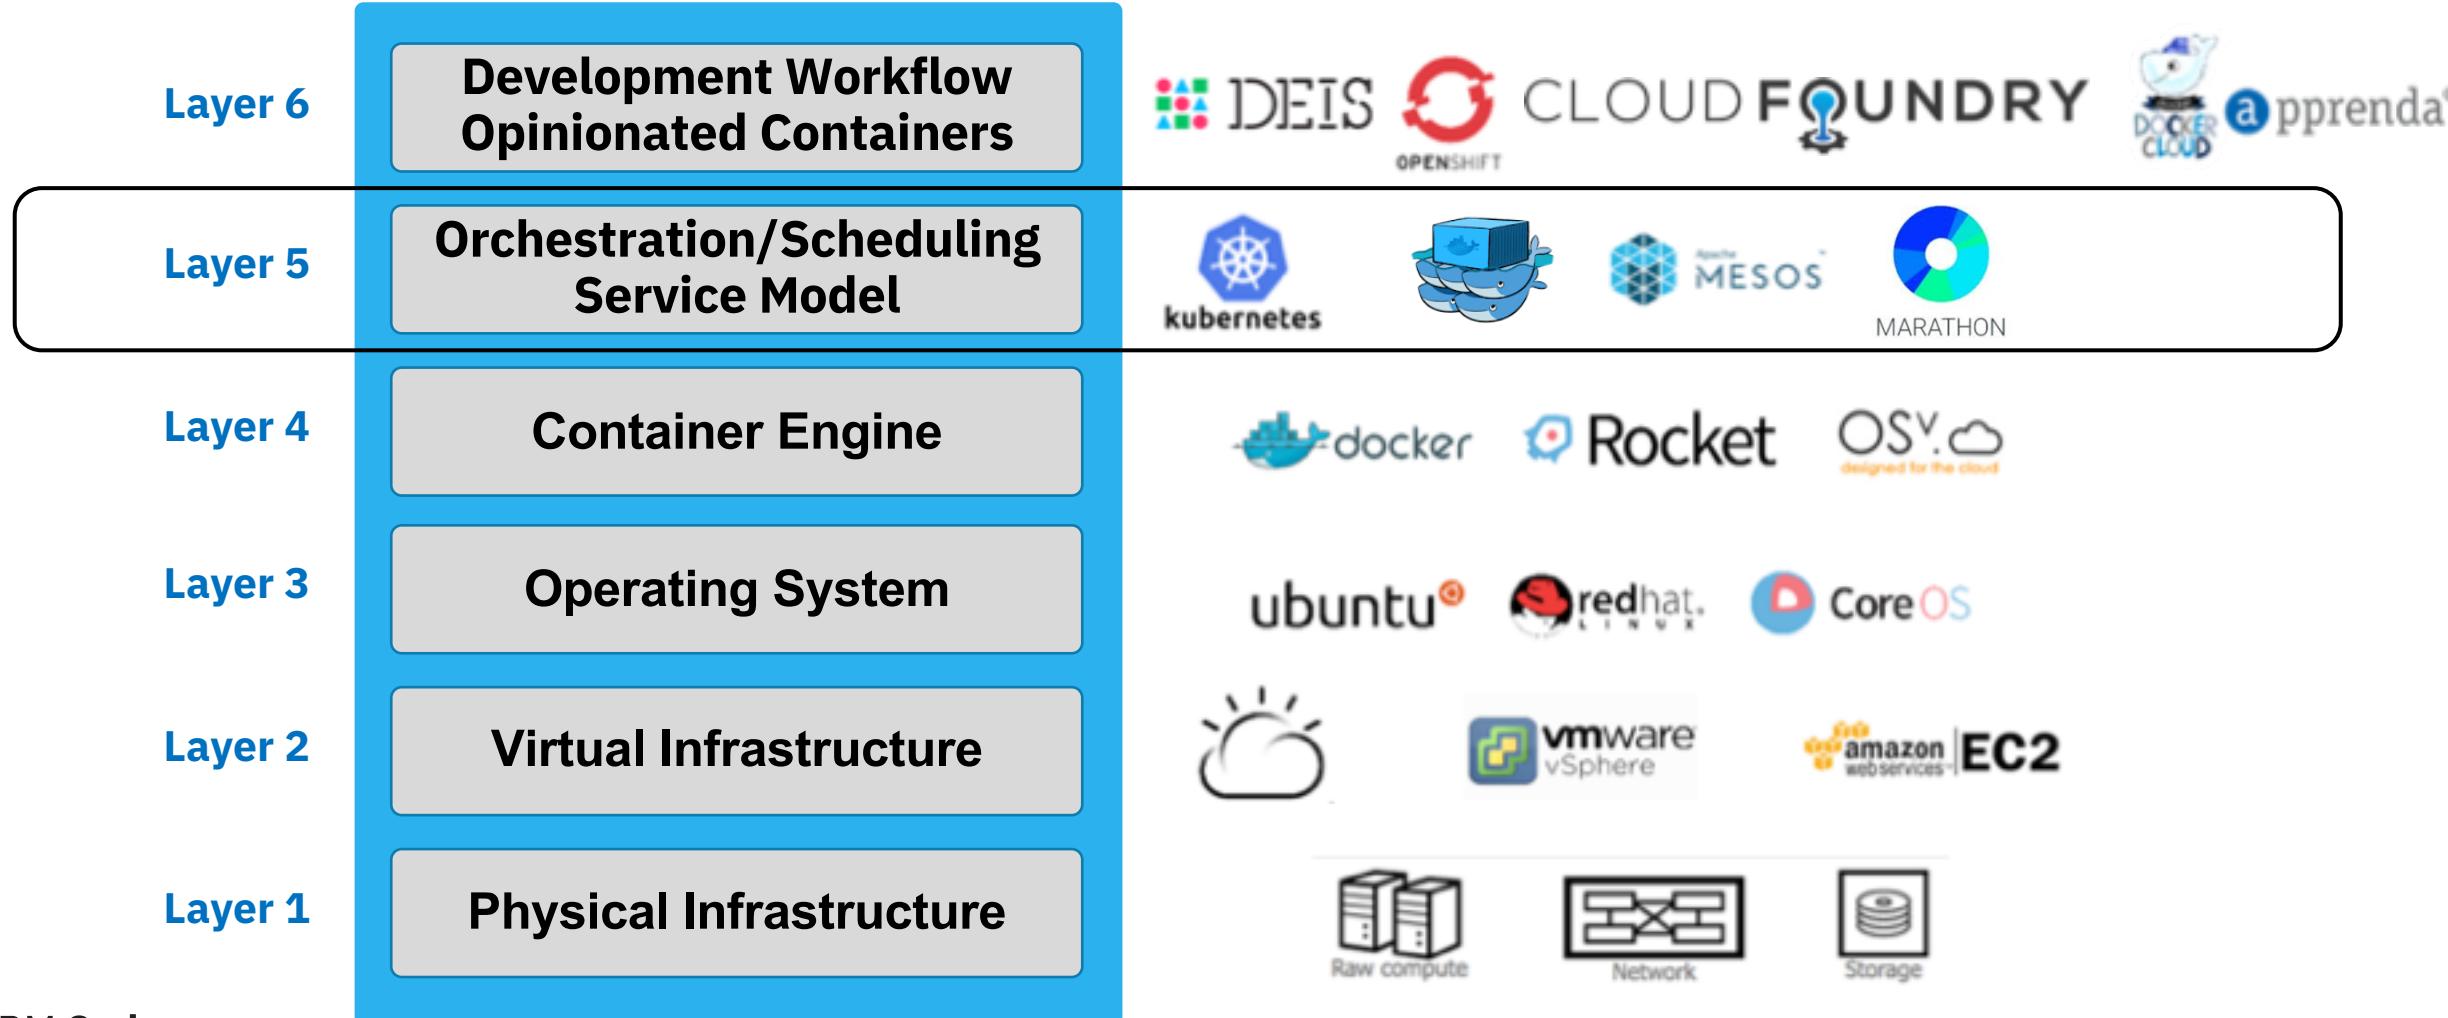
Orchestration

Container Orchestration

Allows users to define how to coordinate the containers in the cloud when the multi-container packaged application is deployed.

- Scheduling
- Cluster management
- Service discovery
- Provisioning
- Monitoring
- Configuration management

Container Ecosystem Layers



Kubernetes



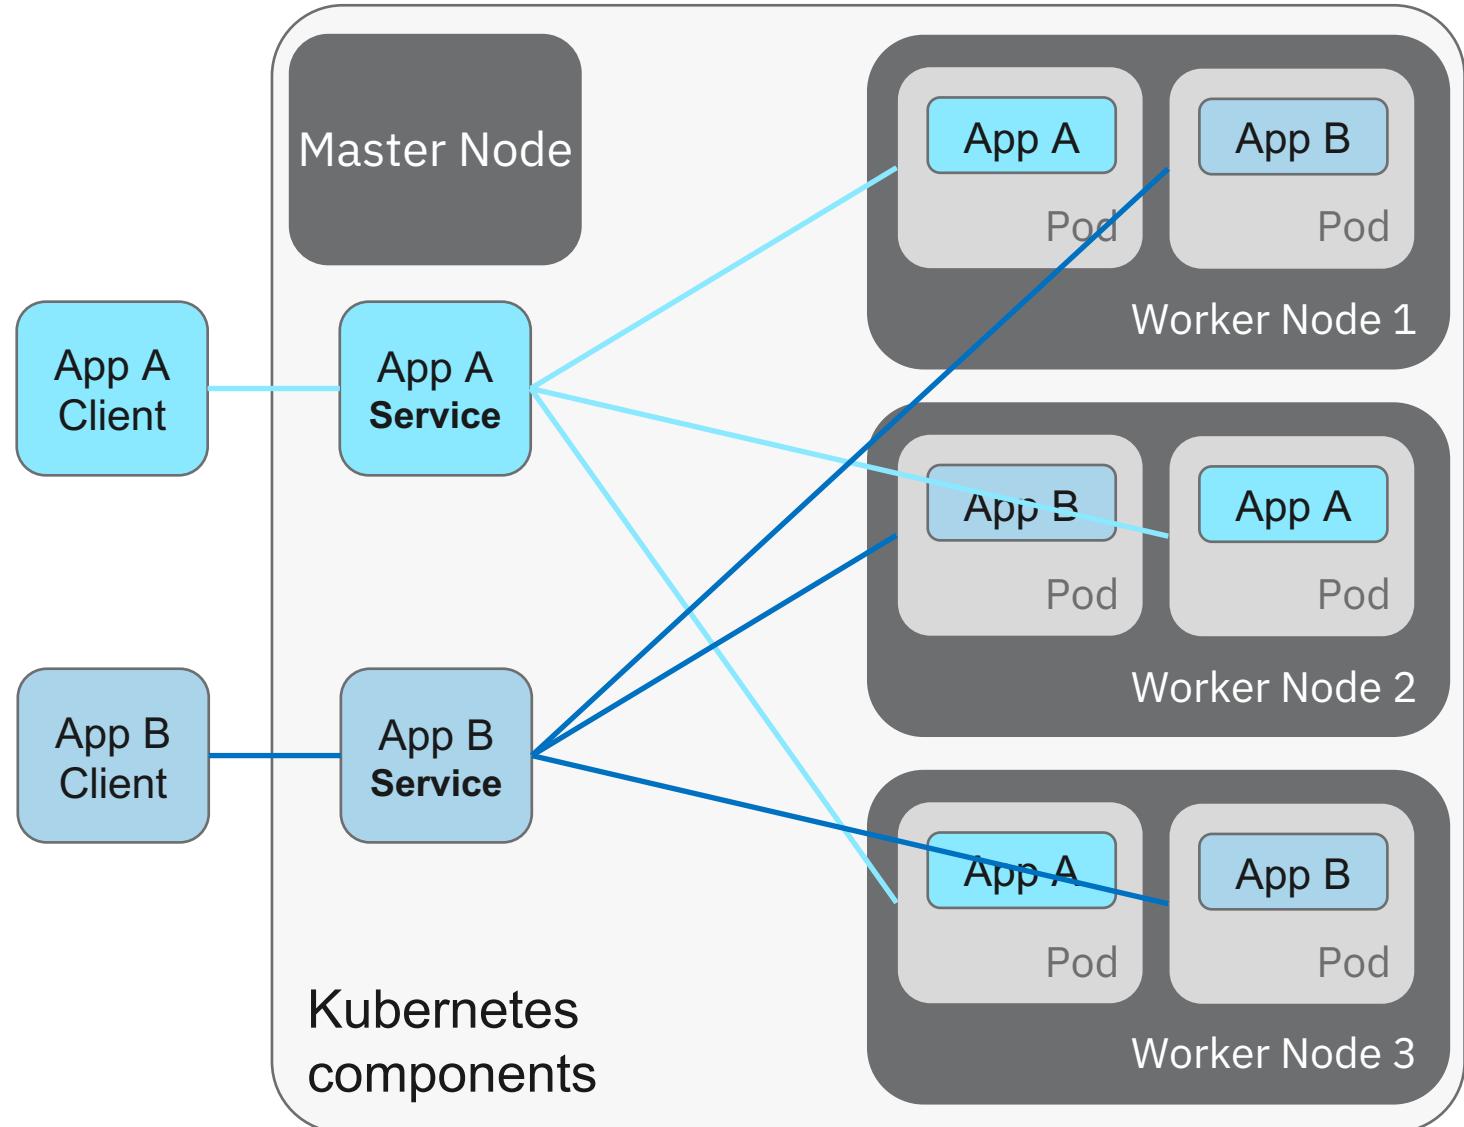
What is Kubernetes?

- Container orchestrator
- Manage applications, **not machines**
- Designed for extensibility
- Open source project managed by the Linux Foundation



Kubernetes Architecture: Workloads

- Container
 - Packaging of an app
- Pod
 - Unit of deployment
- Service
 - Fixed endpoint for 1+ pods



Kubernetes Capabilities

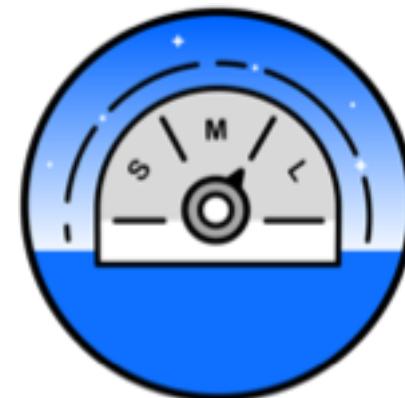




Intelligent Scheduling



Self-healing



Horizontal scaling



Service discovery & load balancing



Automated rollouts and rollbacks



Secret and configuration management

Cluster Management Capabilities

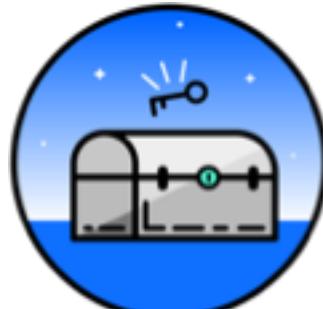




Simplified cluster
management



Design your
own cluster



Container security
& isolation



Extend with
IBM Cloud & Watson



Native open-source
experience



Integrated
operational tools



Design Your Own Cluster

- Tunable capacity
- Select between shared and dedicated compute using virtual server instances
- Bare metal worker nodes
- Edge nodes
- Configurable networking and storage
- Integrated VPN in-cluster providing IPSec tunnels



Container Security & Isolation

- Provides Docker image and running container scanning to detect vulnerabilities and configuration weaknesses with Vulnerability Advisor
- Image security enforcement
- Store your images securely in your hosted private registry
- Docker image signing
- Automatic encryption of secrets and volumes
- Isolated networking and storage
- Default LUKS encryption of /var/lib/docker
 - Every worker node in each cluster has a unique encryption key

Multizone Kubernetes

In June 2018 IBM announced the availability of multizone clusters. Now within the same IBM Cloud region, you can deploy a Kubernetes cluster with worker nodes distributed across different zones (i.e., physical datacenters) ensuring your apps and clusters remain highly available.

The screenshot shows the 'Create new cluster' page on the IBM Cloud website. At the top, there's a navigation bar with the IBM Cloud logo and a 'View All' link. Below it, the title 'Create new cluster' is displayed. A descriptive text explains that the user is provisioning a cluster of hosts called worker nodes to deploy and manage highly-available apps. There are two tabs at the bottom left: 'Docs' and 'Terms'. On the right, the 'Region' is set to 'US South'. Under 'Cluster type', there are two options: 'Free' and 'Standard'. The 'Standard' option is selected, indicated by a checked checkbox icon. A brief description for the Standard cluster states: 'Ready for production? Create a fully-customizable cluster with your choice of hardware isolation.' and 'Starting from \$0.11 hourly'. Below this, the 'Location' section includes an 'Availability' dropdown with 'Single Zone' and 'Multizone' options, where 'Multizone' is selected. The 'Zones' section lists three zones: 'dal10', 'dal12', and 'dal13', each associated with a Private VLAN (e.g., '2327387-2261-bcr01a.dal10') and a Public VLAN (e.g., '2327385-1826-fcr01a.dal10').

IBM Cloud Kubernetes – try Kubernetes for free!

IBM Cloud

Catalog Docs Support Manage

[View All](#)

Create new cluster

Provision a cluster of hosts, called worker nodes, to deploy and manage highly-available apps.

Region [?](#)

US East

Plan type

Free [?](#)
New to Kubernetes? Create a cluster with 1 worker node to explore the capabilities.
[Free](#)

Pay-As-You-Go plan [?](#)
Create a fully-customizable, production-ready cluster with your choice of hardware
[Starting from \\$0.19 hourly](#)

Cluster details

Cluster Name

Location [?](#)

Kubernetes version [?](#)

Order Summary

Free - 2 CPUs, 4 GB RAM

1 worker node [Free](#)

Total due now: [Free](#) estimated

[Create Cluster](#)

[Cancel](#)

Istio



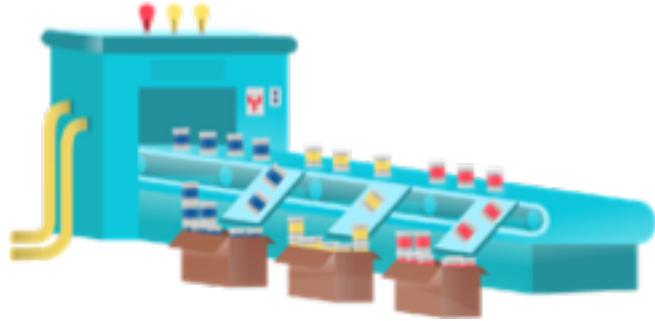
An open platform to connect, manage, and secure microservices

<http://istio.io>

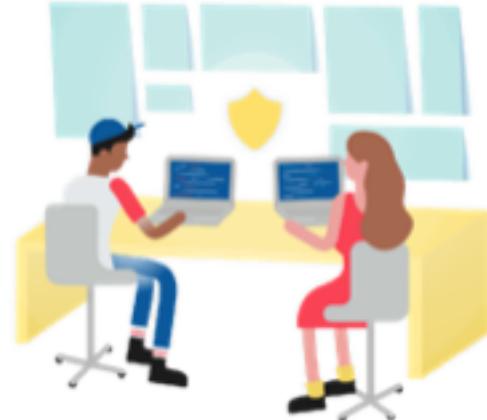
Istio

WHERE ISTIO FINDS ITS PLACE

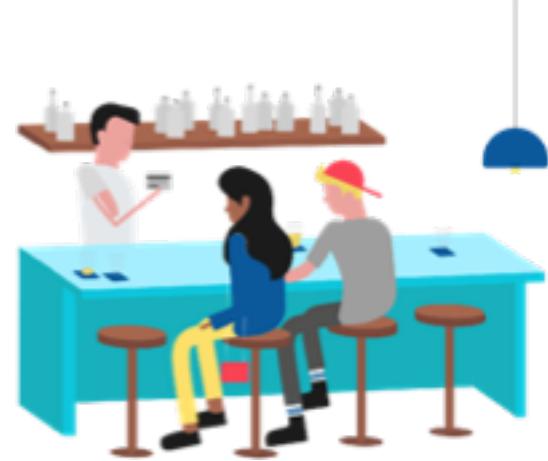
Intelligent
Routing and Load
Balancing



Resiliency across
Languages and
Platforms



Fleet Wide Policy
Enforcement



In-Depth
Telemetry and
Reporting



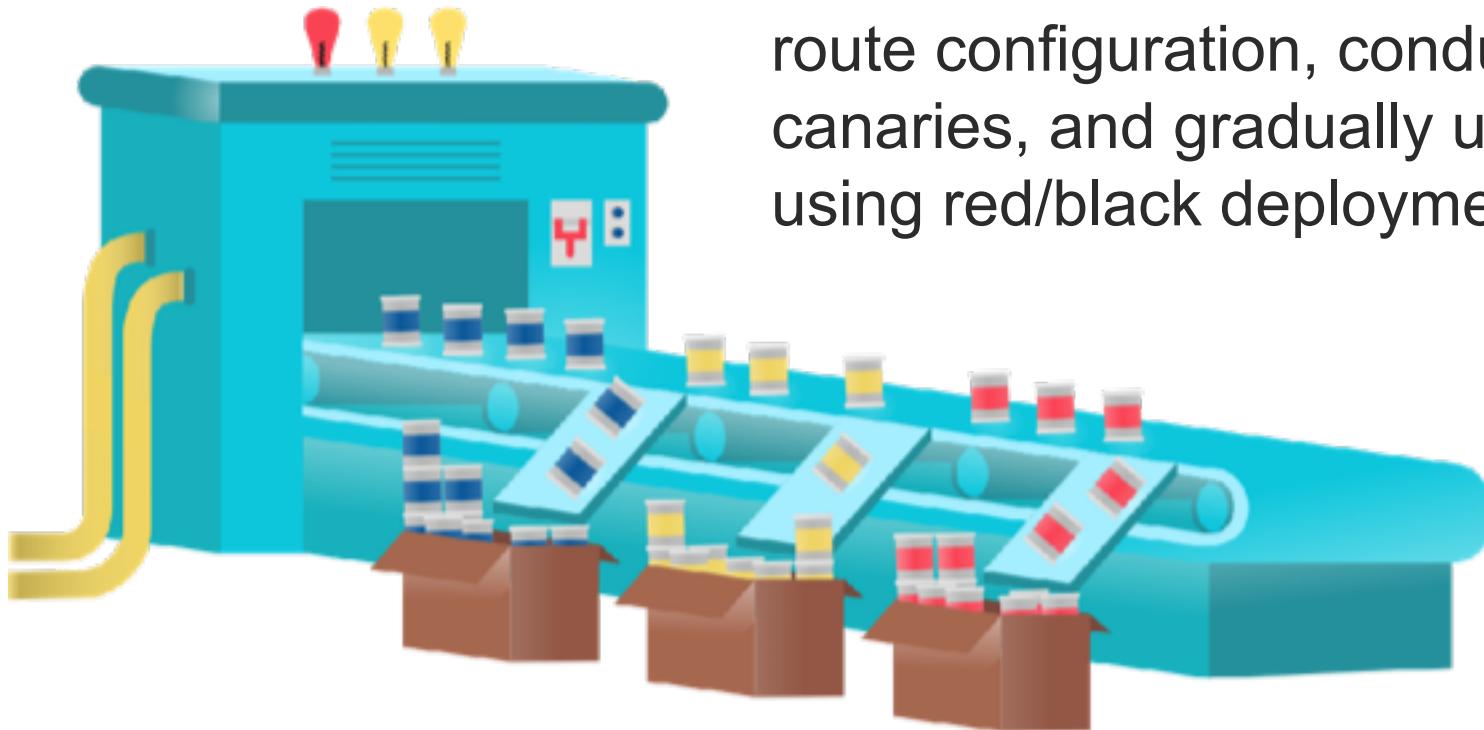
IBM Code

Istio

WHERE ISTIO FINDS ITS PLACE

Intelligent Routing and Load Balancing

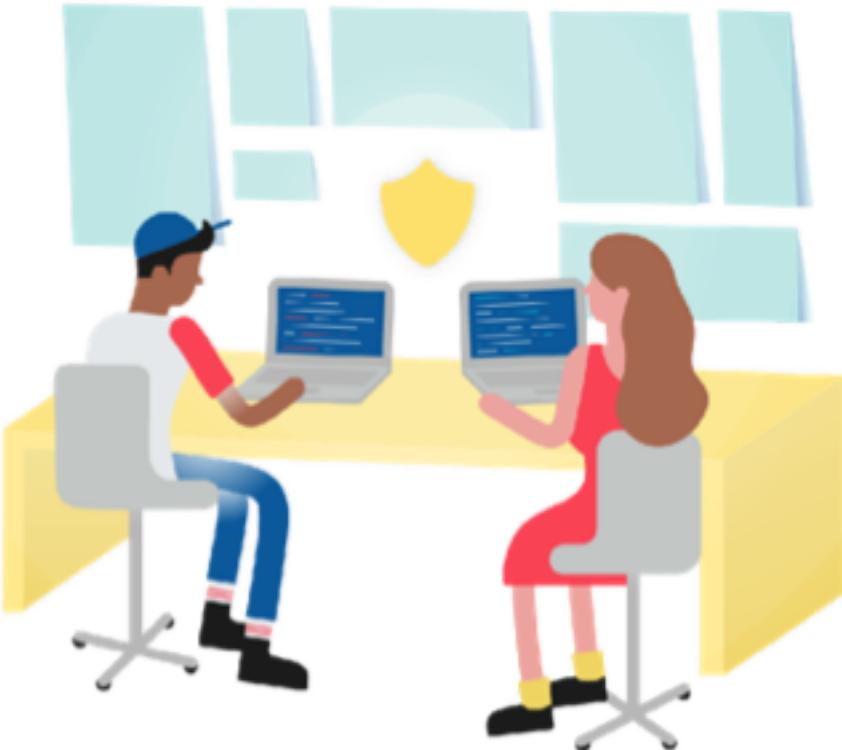
Control traffic between services with dynamic route configuration, conduct A/B tests, release canaries, and gradually upgrade versions using red/black deployments.



Istio

WHERE ISTIO FINDS ITS PLACE

Resiliency across Languages and Platforms



Increase reliability by shielding applications from flaky networks and cascading failures in adverse conditions.



IBM Code

Istio

WHERE ISTIO FINDS ITS PLACE

Fleet Wide Policy Enforcement



Apply organizational policy to the interaction between services, ensure access policies are enforced and resources are fairly distributed among consumers.

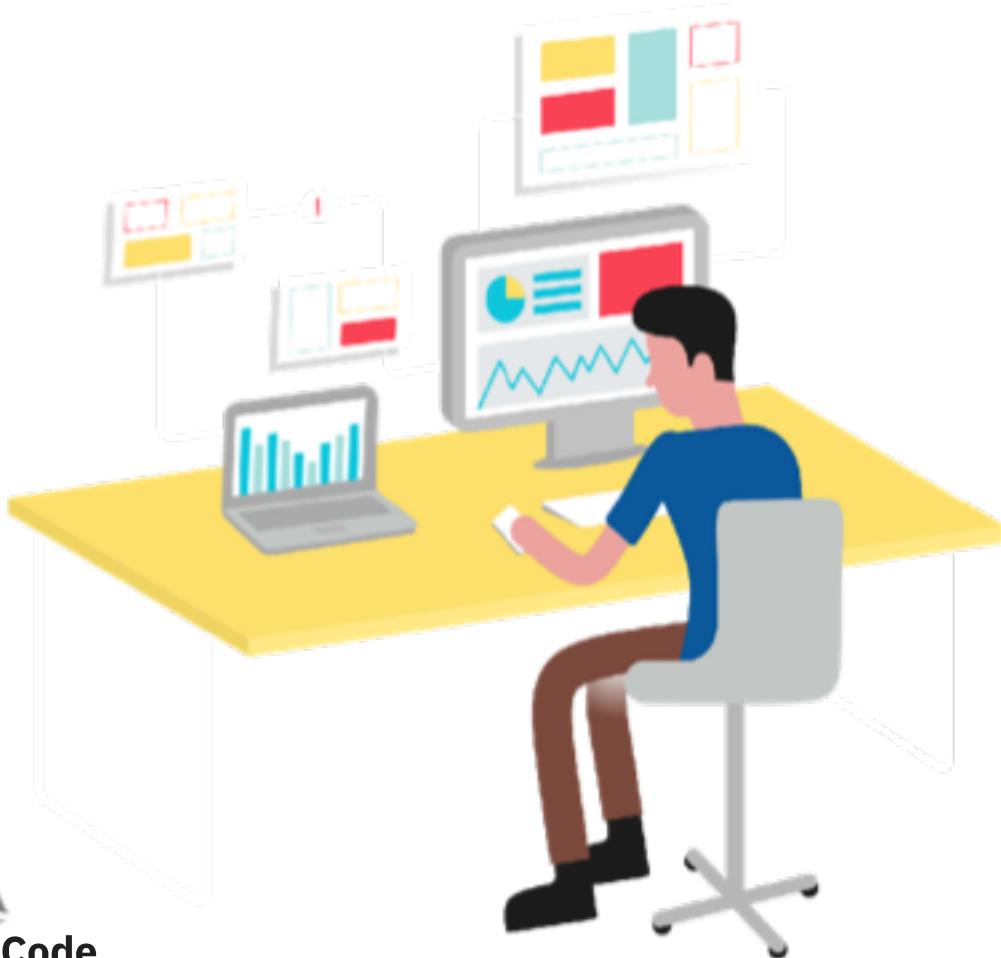


IBM Code

Istio

WHERE ISTIO FINDS ITS PLACE

In-Depth Telemetry and Reporting

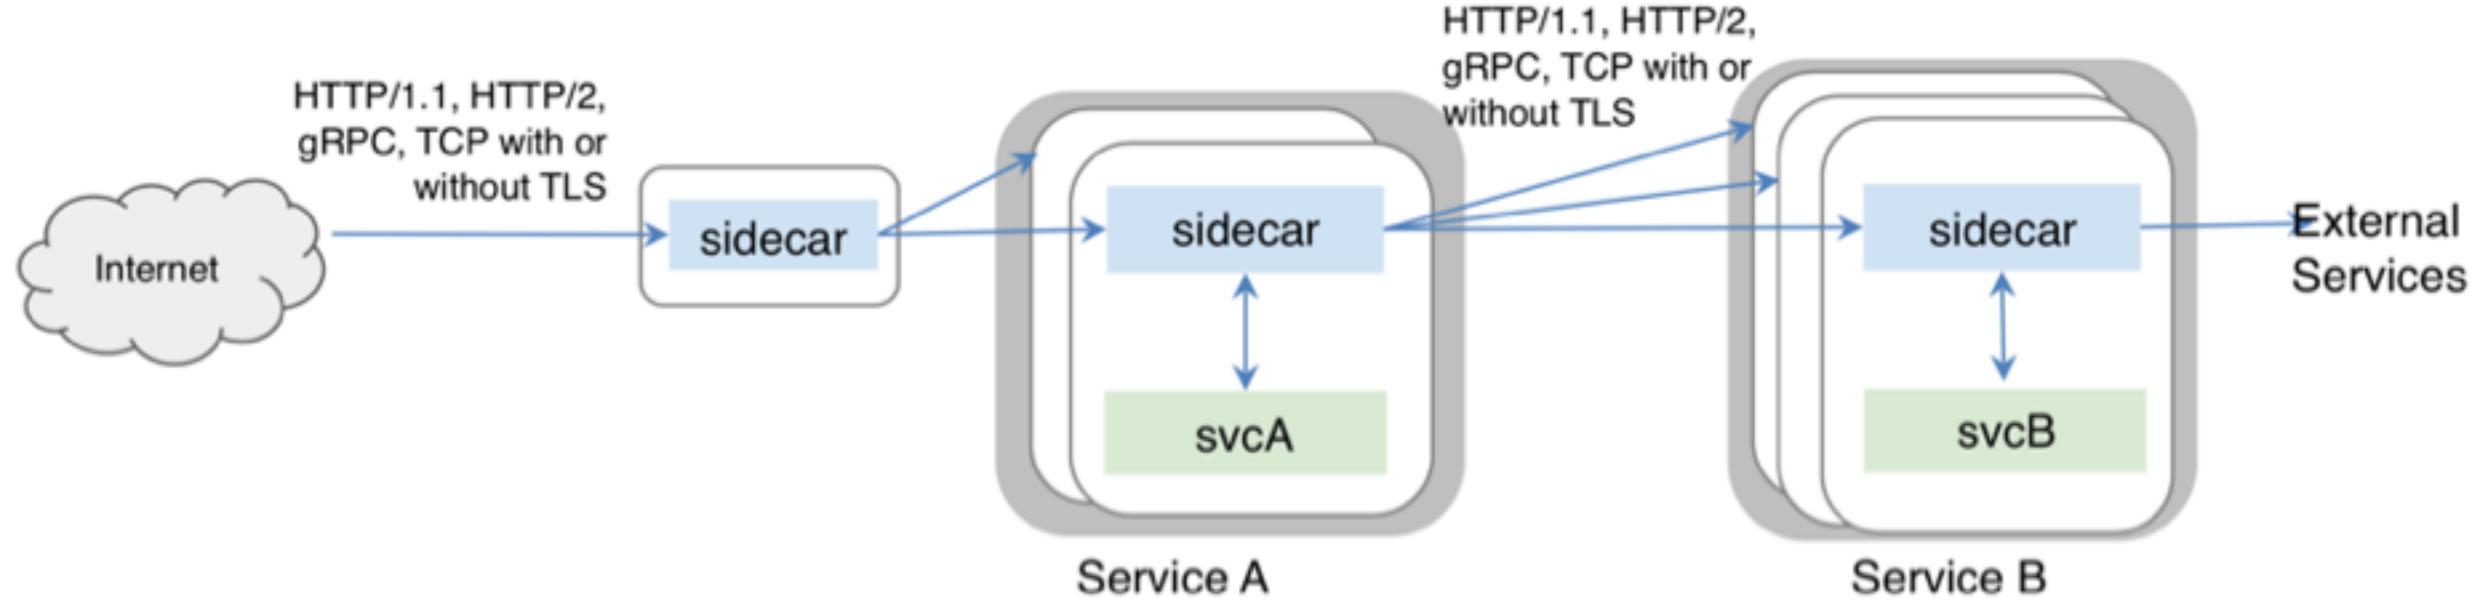


Understand the dependencies between services, the nature and flow of traffic between them and quickly identify issues with distributed tracing.



IBM Code

Istio – original concept - Sidecar



Outbound features:

- Service authentication
- Load balancing
- Retry and circuit breaker
- Fine-grained routing
- Telemetry
- Request Tracing
- Fault Injection

Inbound features:

- Service authentication
- Authorization
- Rate limits
- Load shedding
- Telemetry
- Request Tracing
- Fault Injection



Envoy by Lyft – sidecar of choice

- A C++ based L4/L7 proxy
- Low memory footprint
- Battle-tested @ Lyft
 - 100+ services
 - 10,000+ VMs
 - 2M req/s



Goodies:

- HTTP/2 & gRPC
- Zone-aware load balancing w/ failover
- Health checks, circuit breakers, timeouts, retry budgets
- No hot reloads - API driven config updates

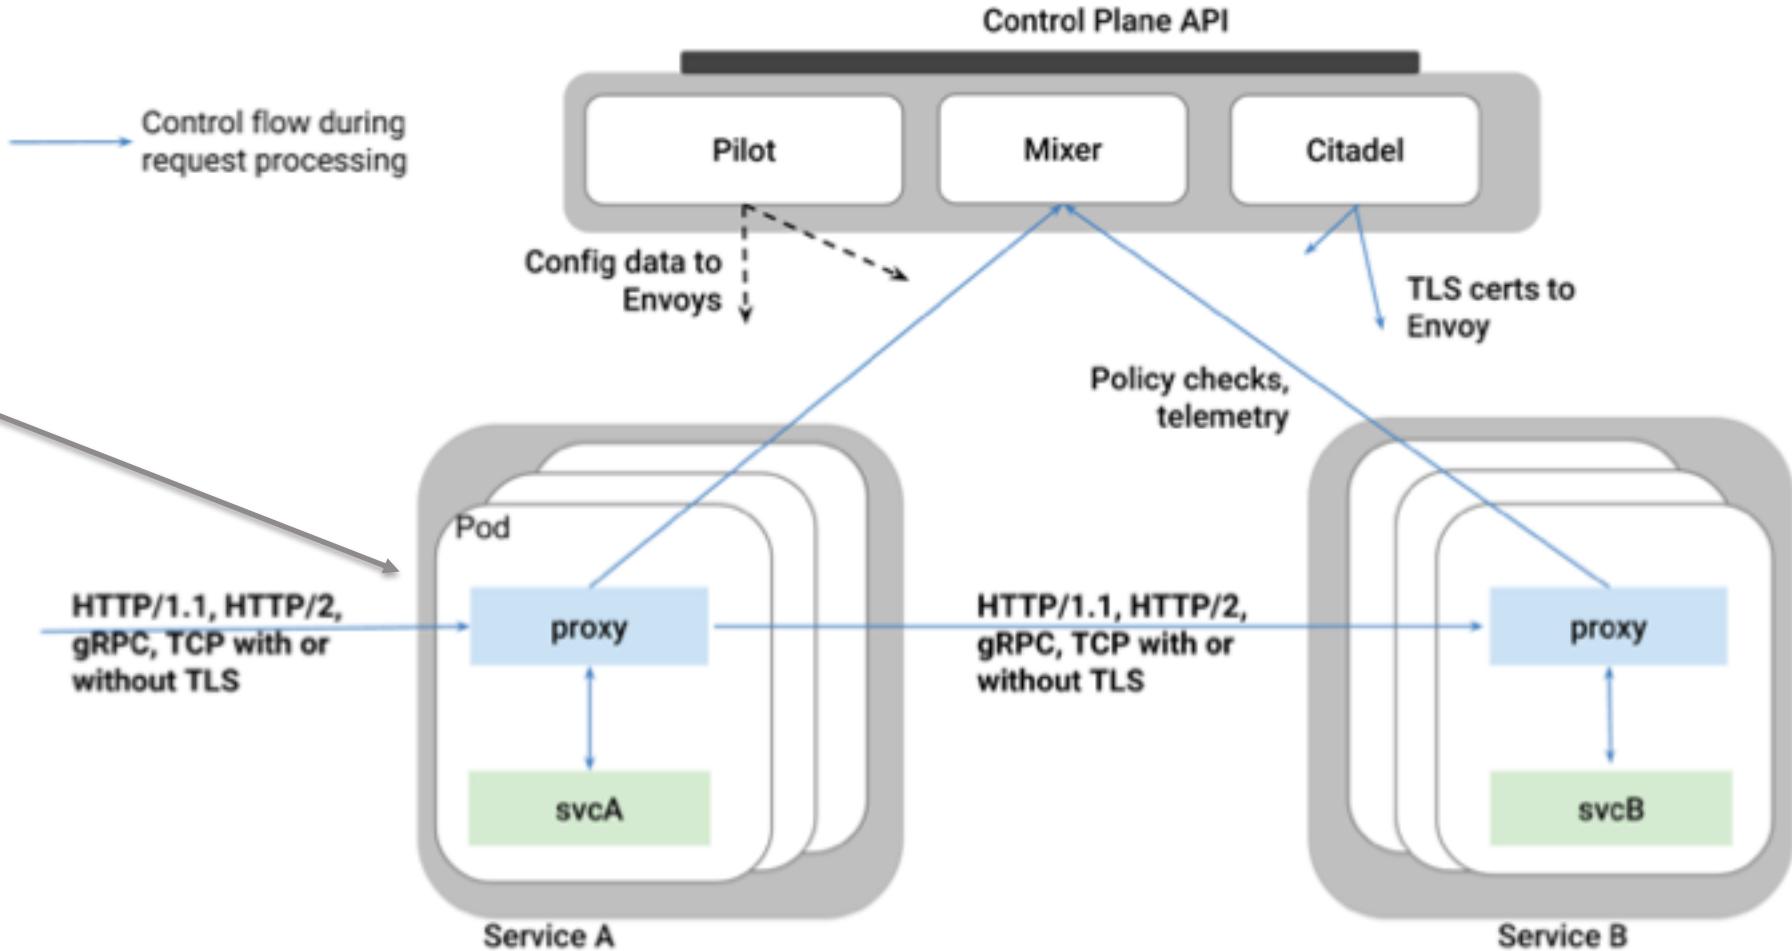
Istio's contributions:

- Transparent proxying w/ SO_ORIGINAL_DST
- Traffic routing and splitting
- Request tracing using Zipkin
- Fault injection



Istio Architecture

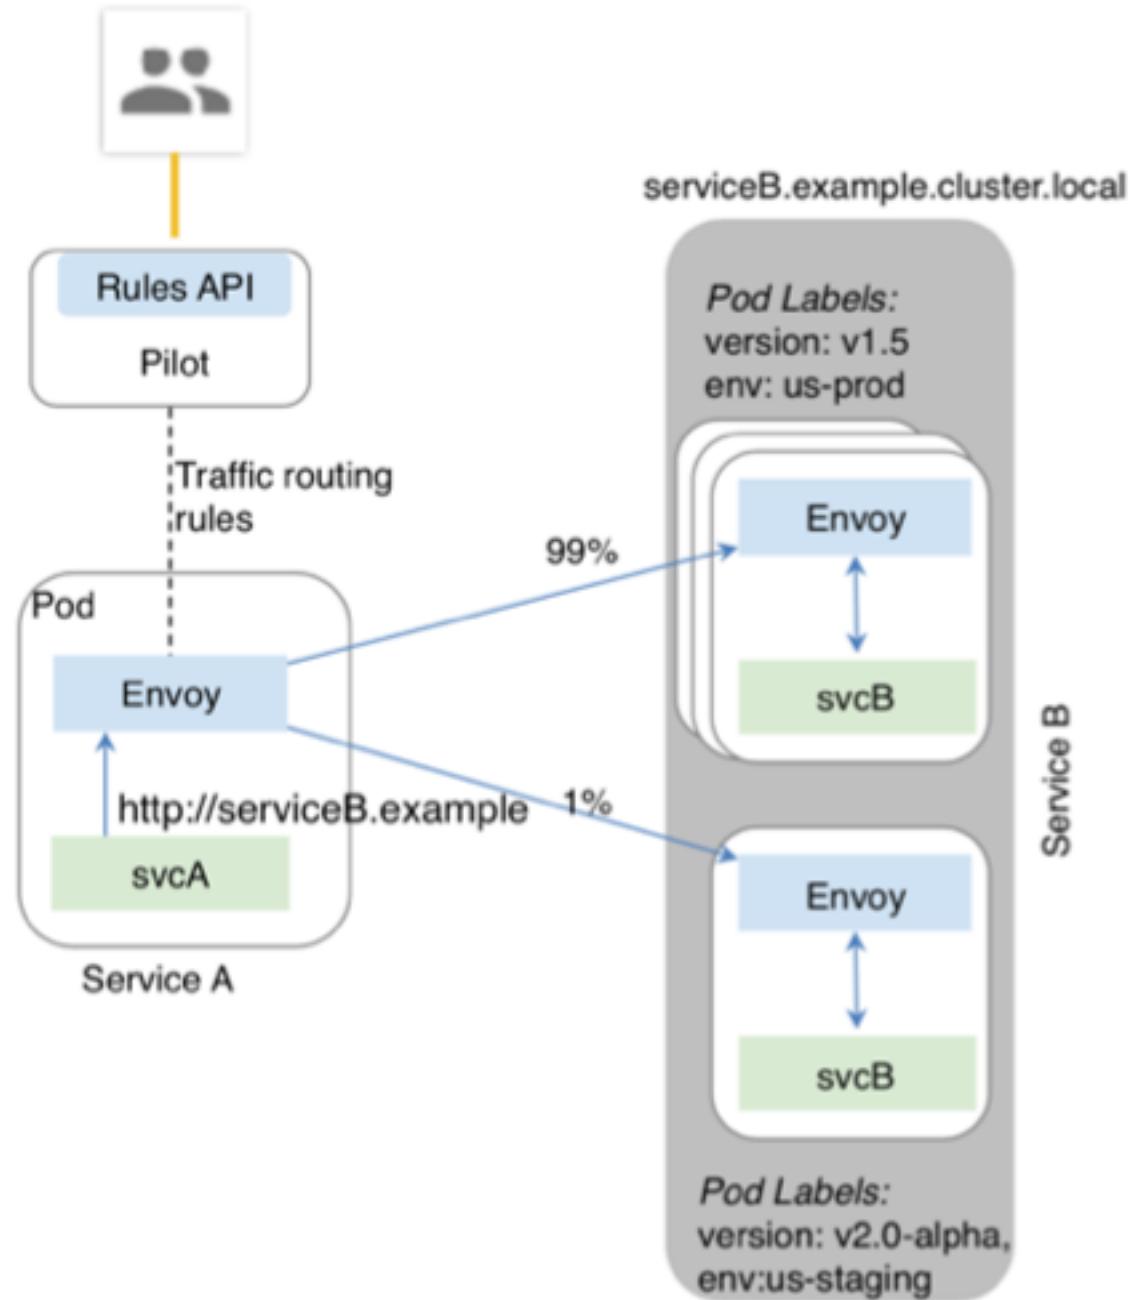
Traffic is transparently intercepted and proxied. An App is unaware of Envoy's presence



Traffic splitting

```
// A simple traffic splitting rule

destination: serviceB.example.cluster.local
match:
  source: serviceA.example.cluster.local
route:
- tags:
    version: v1.5
    env: us-prod
    weight: 99
- tags:
    version: v2.0-alpha
    env: us-staging
    weight: 1
```



Traffic control is decoupled from infrastructure scaling

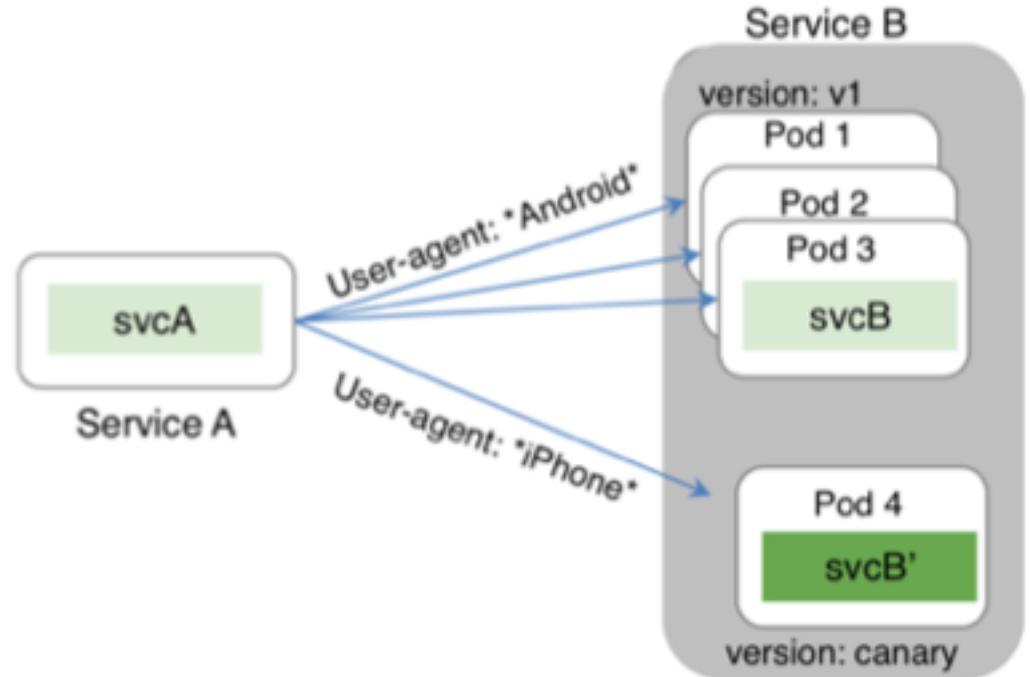
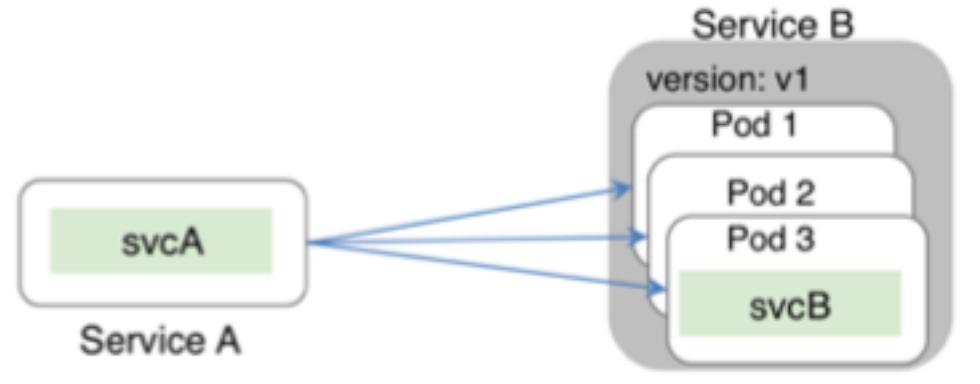


Traffic steering

```
// Content-based traffic steering rule

destination: serviceB.example.cluster.local
match:
  httpHeaders:
    user-agent:
      regex: ^(.*)?(iPhone)(;.*?)$
precedence: 2
route:
- tags:
  version: canary
```

Content-based traffic steering



Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers

destination: serviceB.example.cluster.local
policy:
- tags:
  version: v1
  circuitBreaker:
    simpleCb:
      maxConnections: 100
      httpMaxRequests: 1000
      httpMaxRequestsPerConnection: 10
      httpConsecutiveErrors: 7
      sleepWindow: 15m
      httpDetectionInterval: 5m
```

Resilience features

- Timeouts
- Retries with timeout budget
- Circuit breakers
- Health checks
- AZ-aware load balancing w/ automatic failover
- Control connection pool size and request load
- Systematic fault injection

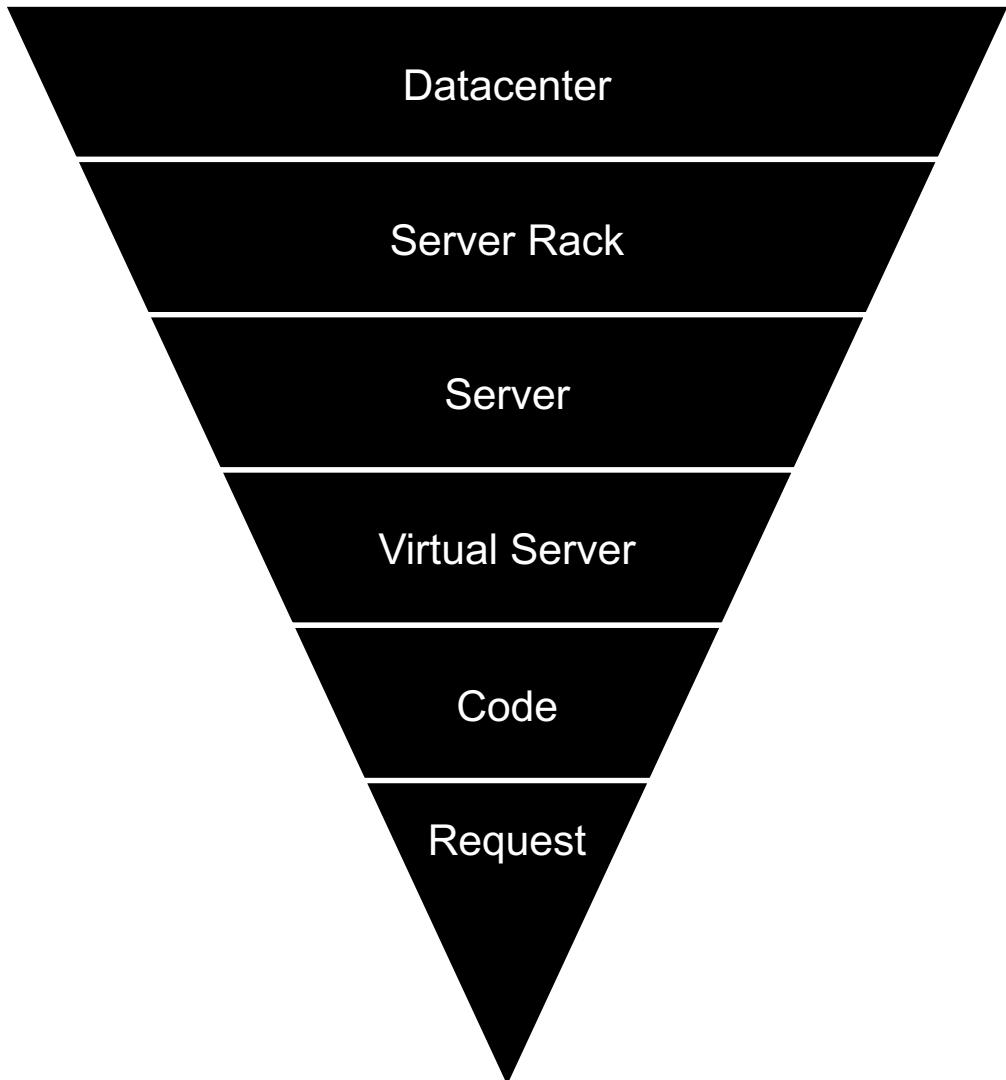


Sometimes things don't go as planned...



Many possibilities for failures

- Failures **WILL** happen
- Their root causes can be complex and affect different number of users
- Reduce the likelihood of correlated failures
- Recover quickly
- Proactive detection of problems before they happen

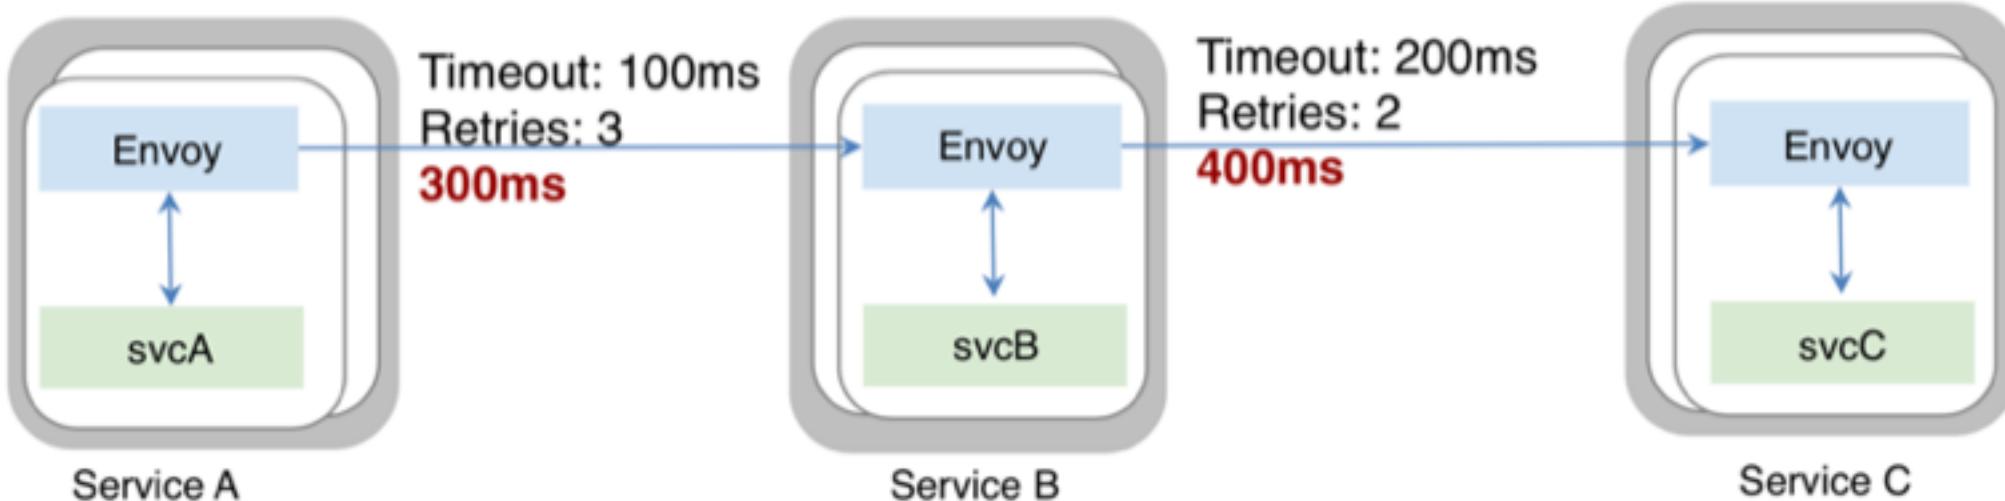


Resiliency testing

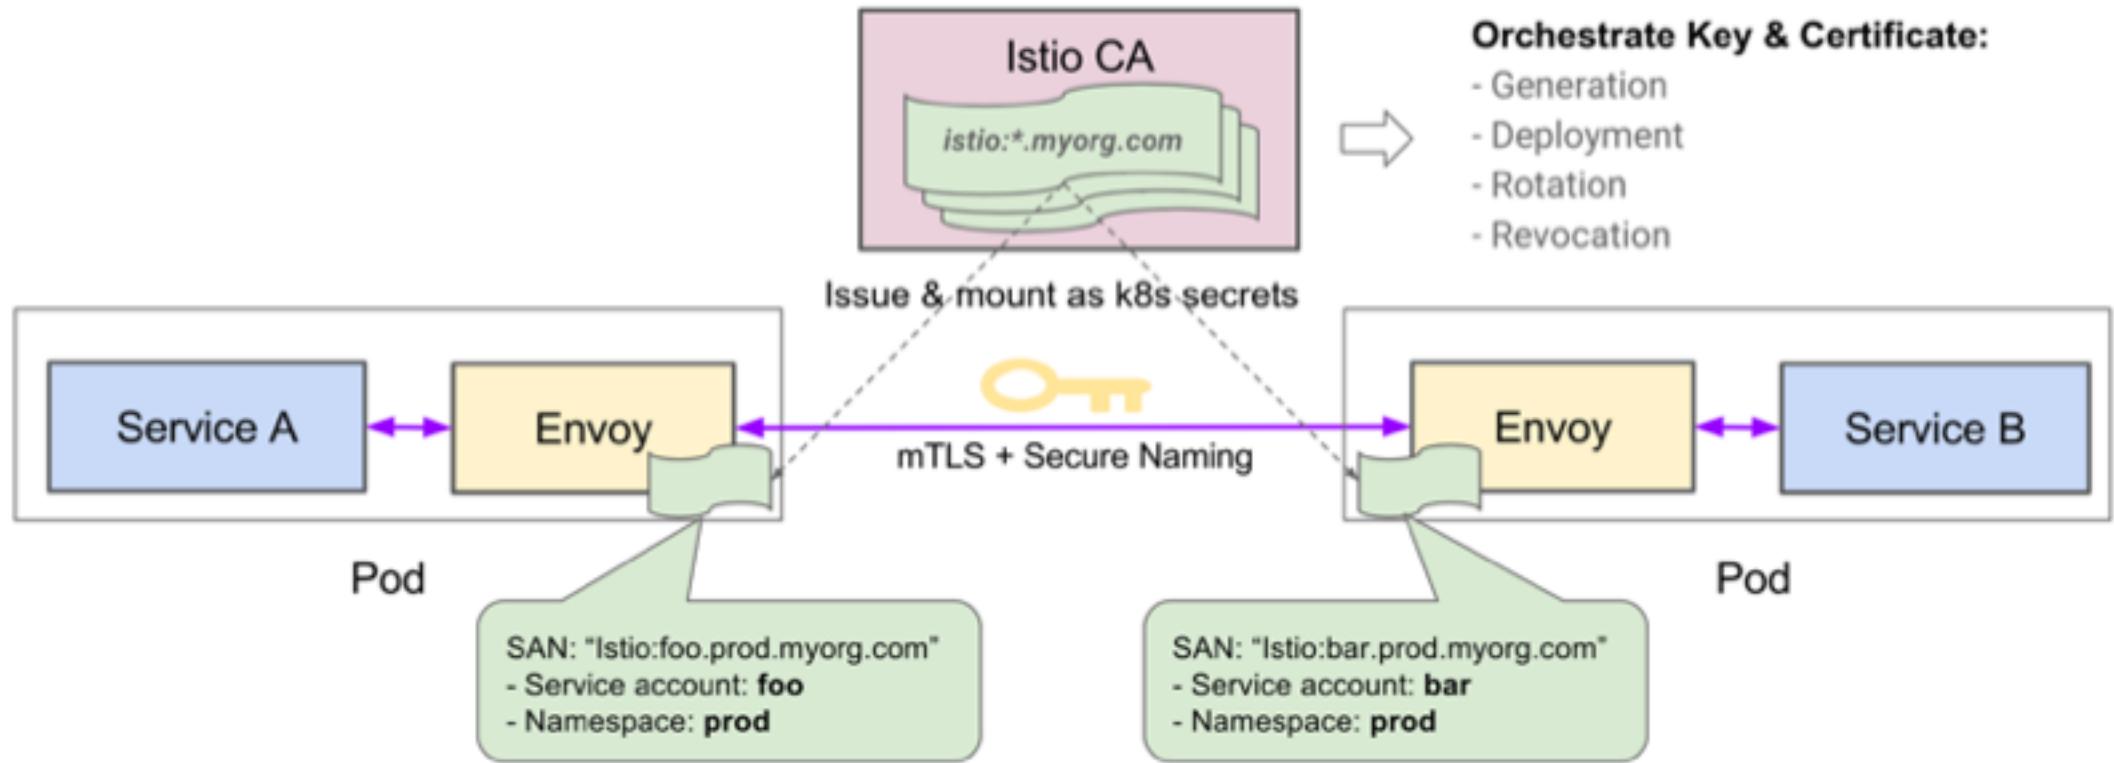
Systematic fault injection to identify weaknesses in failure recovery policies:

HTTP/gRPC error codes

Delay injection



Istio-Auth

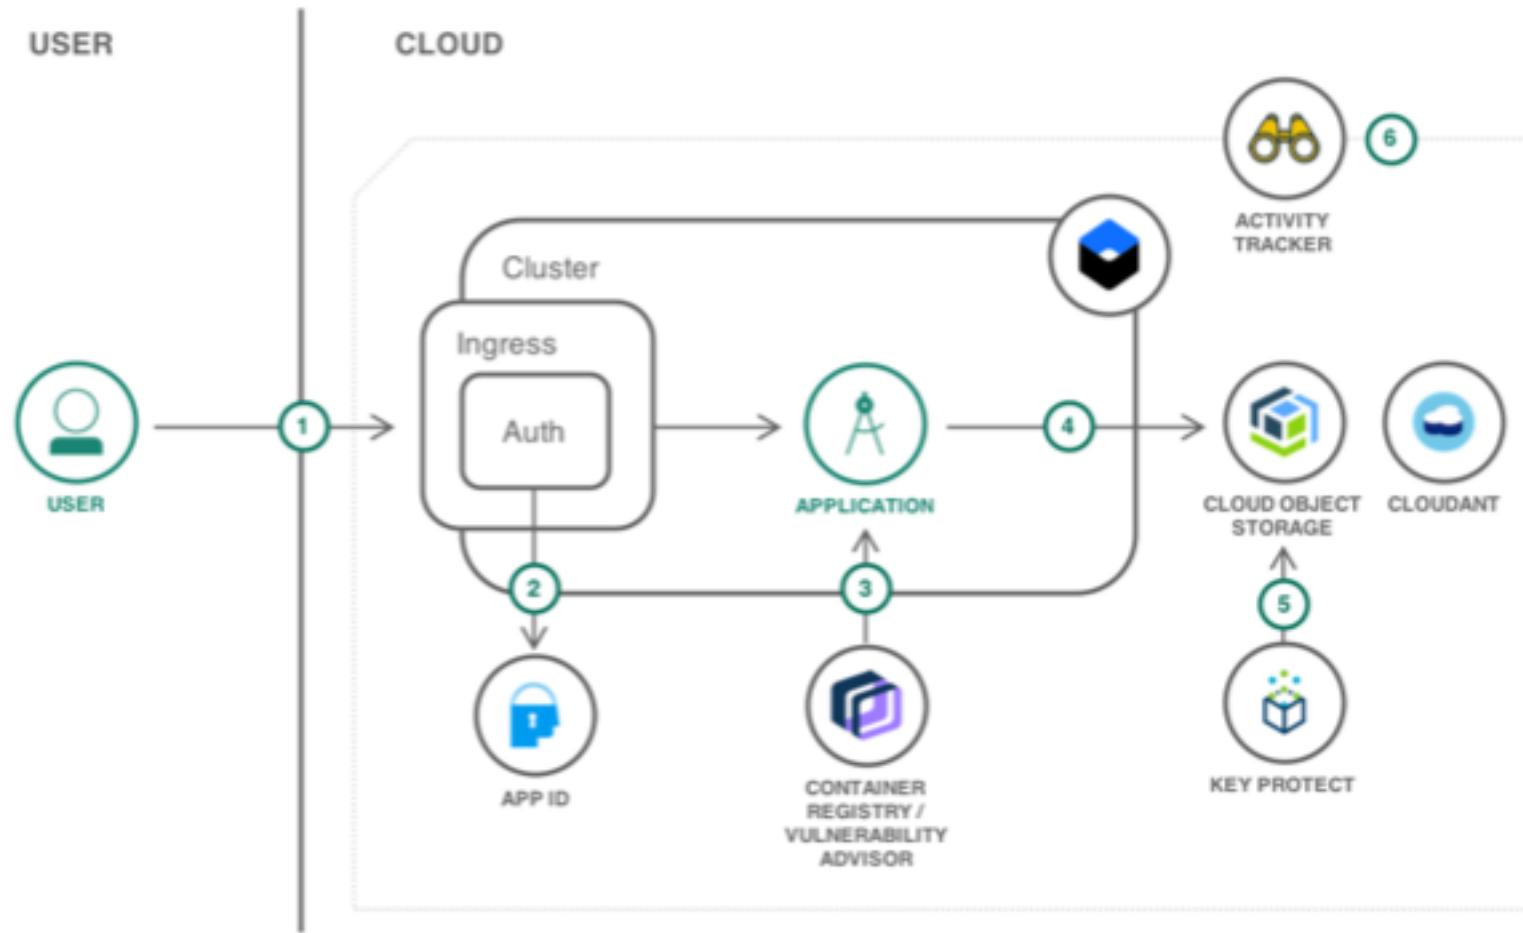


Conclusion

Conclusion

- Why containers?
- Docker containers
- Container orchestration
- Istio – beyond smart routing and resiliency
- Get started with Istio and IBM Cloud Kubernetes Service
 - <https://developer.ibm.com/courses/all/get-started-istio-ibm-cloud-container-service>

The architecture



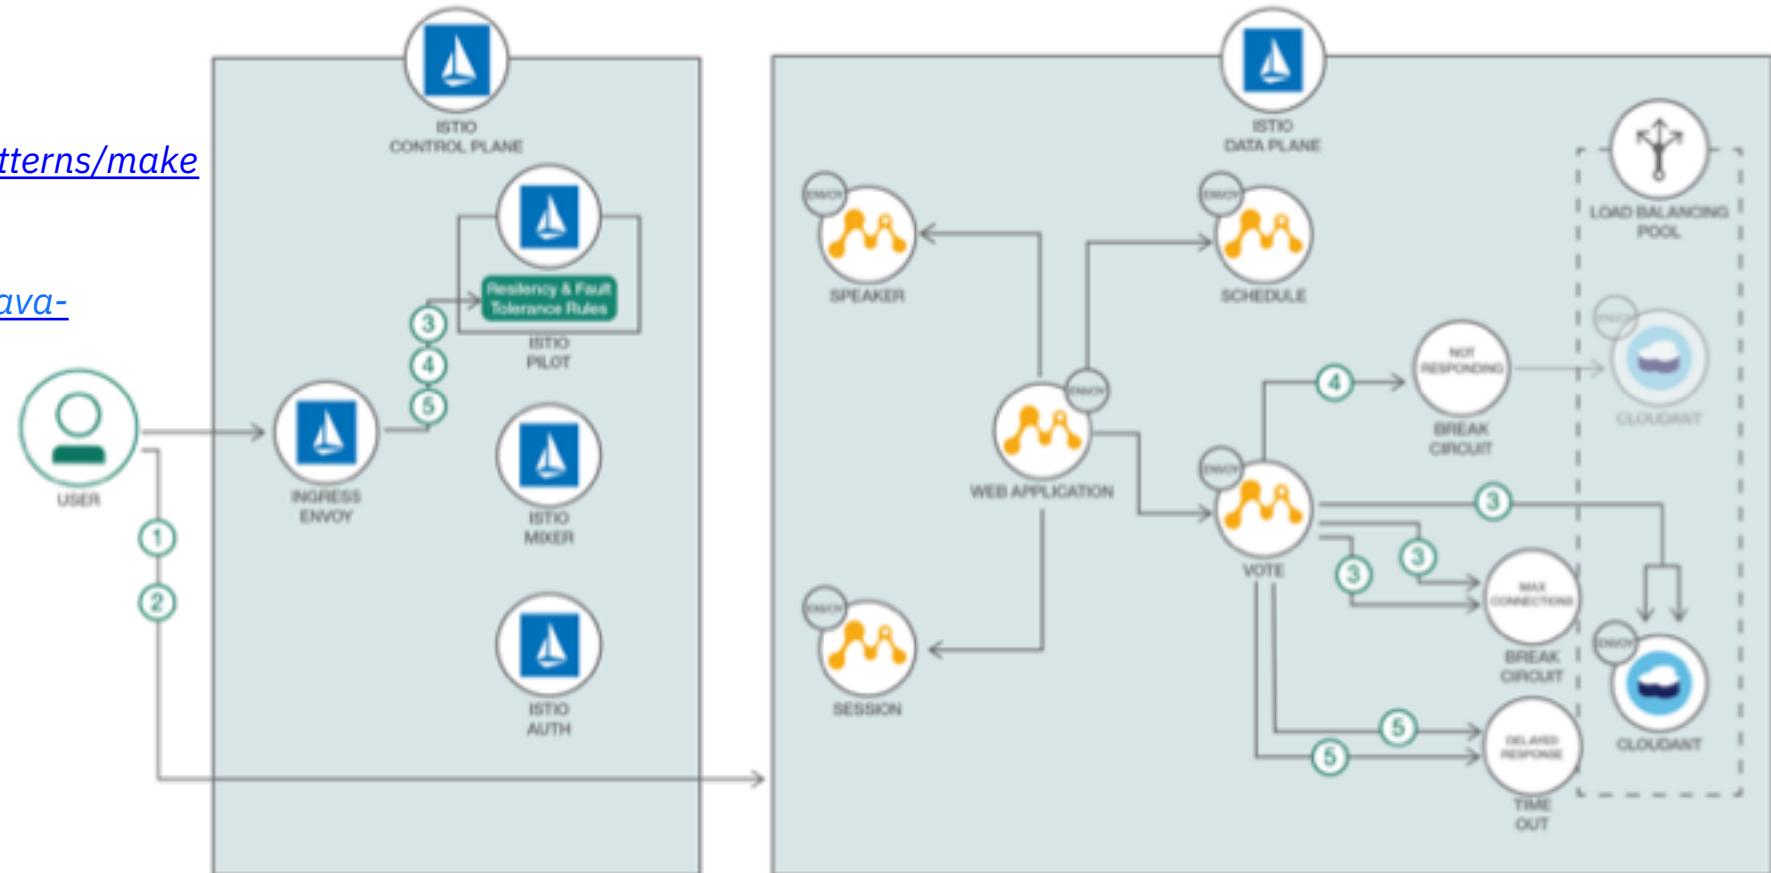
Java Based Code Example – aka pattern

Enable your Java microservices with advanced resiliency features leveraging Istio

Developer Works

Code: <https://developer.ibm.com/code/patterns/make-java-microservices-resilient-with-istio/>

Github: <https://github.com/IBM/resilient-java-microservices-with-istio>



Resources

- Docker tutorial - <https://docs.docker.com/get-started/>
- Kubernetes tutorial - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
- The Evolution of Linux Containers and Their Future
 - <https://dzone.com/articles/evolution-of-linux-containers-future>
- Introduction to container orchestration
 - <https://www.exoscale.ch/syslog/2016/07/26/container-orch/>
- TNS Research: The Present State of Container Orchestration
 - <https://thenewstack.io/tns-research-present-state-container-orchestration/>
- Large-scale cluster management at Google with Borg
 - <https://research.google.com/pubs/pub43438.html>
- Call for code: <http://developer.ibm.com/callforcode>
- The other SF City events: <https://developer.ibm.com/code/community/cities/san-francisco/>

Thank you

-  twitter.com/blumareks
-  github.com/blumareks
-  developer.ibm.com/code

